


# Microservice Communication

*Rob Vettor*



# Microservices Communication

A large challenge  
moving to a  
microservice-based  
architecture is  
communication




We'll discuss...

Front-End  
Client Communication

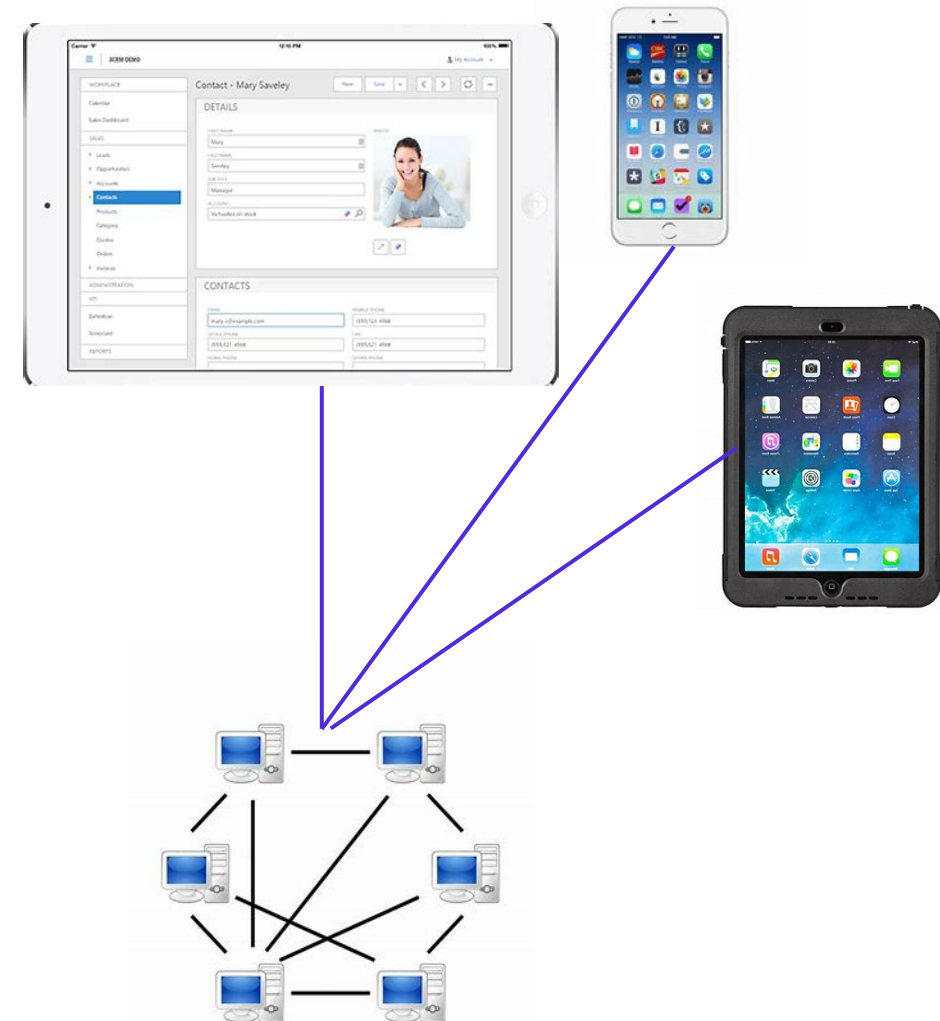
Backend  
Service Communication

# Front-end Client Communication



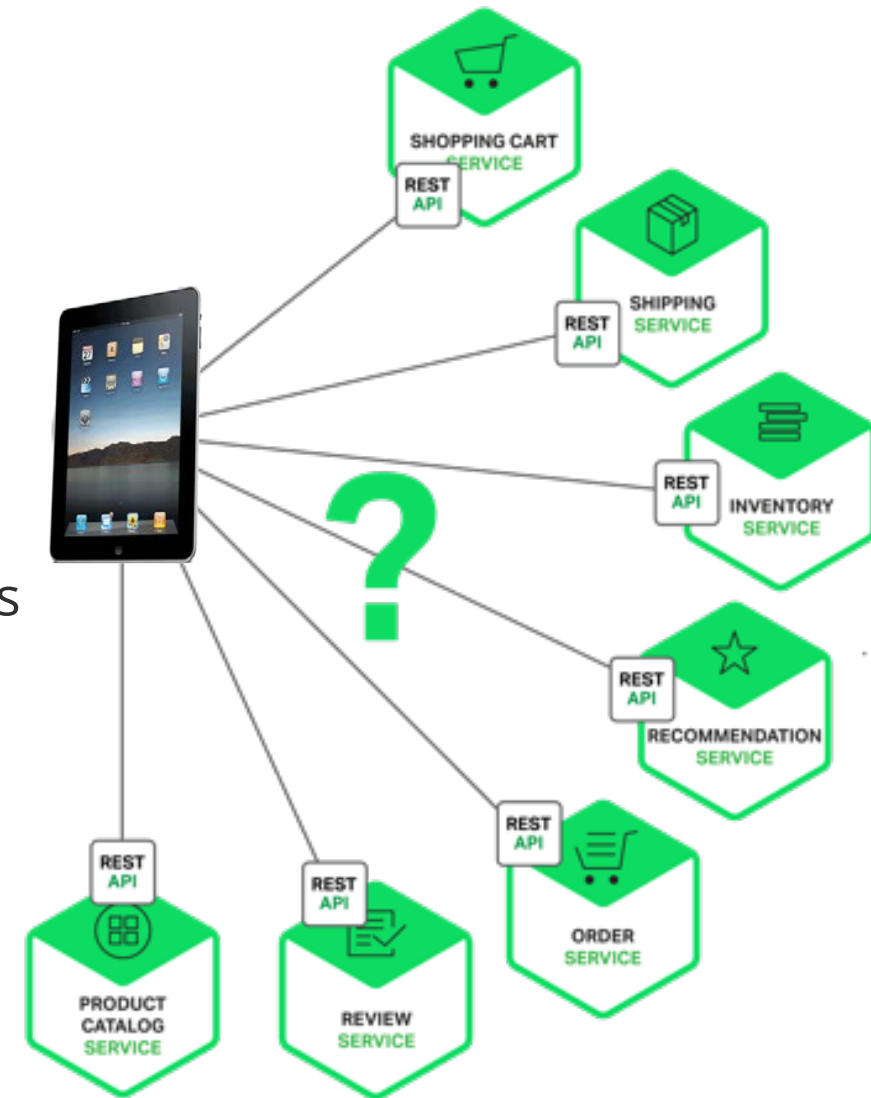
# Front-End Clients

- Distributed applications typically have *decoupled* front-end clients
- What communication patterns are in play?
  - Direct client communication
  - API Gateway pattern
- Both commonly implemented with REST (HTTP), but gRPC is gaining popularity



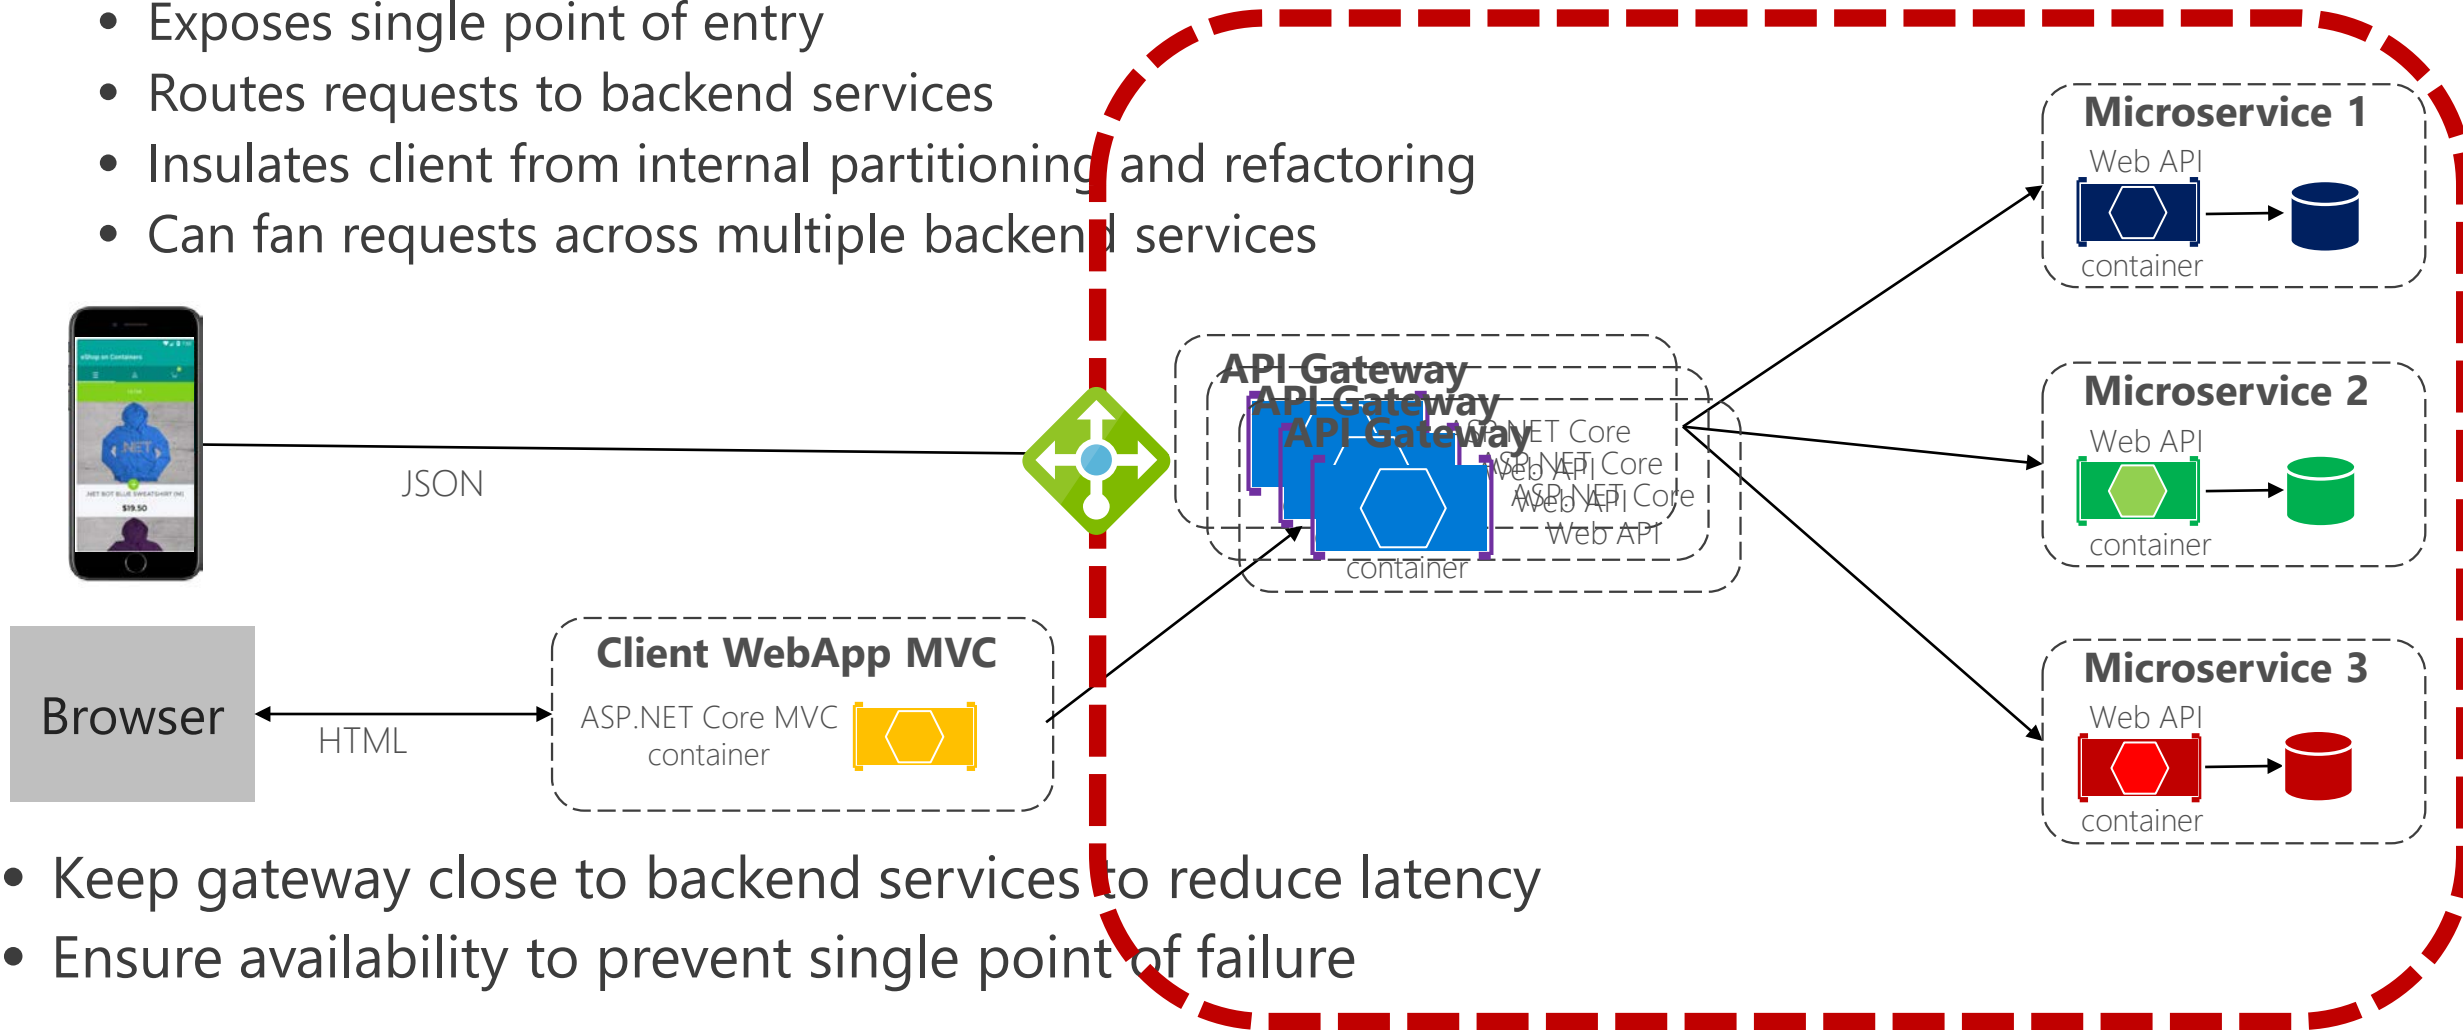
# Direct Client Communication

- Client communicates directly with each service
- Simple, but far from ideal...
  - Client is tightly-coupled to backend core services
  - Client susceptible to service refactoring/partitioning
  - Client can become chatty – orchestrating multiple calls
  - Client becomes complex – often contains business logic
  - Each service must support full set of cross-cutting concerns
- Security concerns: Direct access exposes backend services and widens attack surface



# API Gateway Pattern

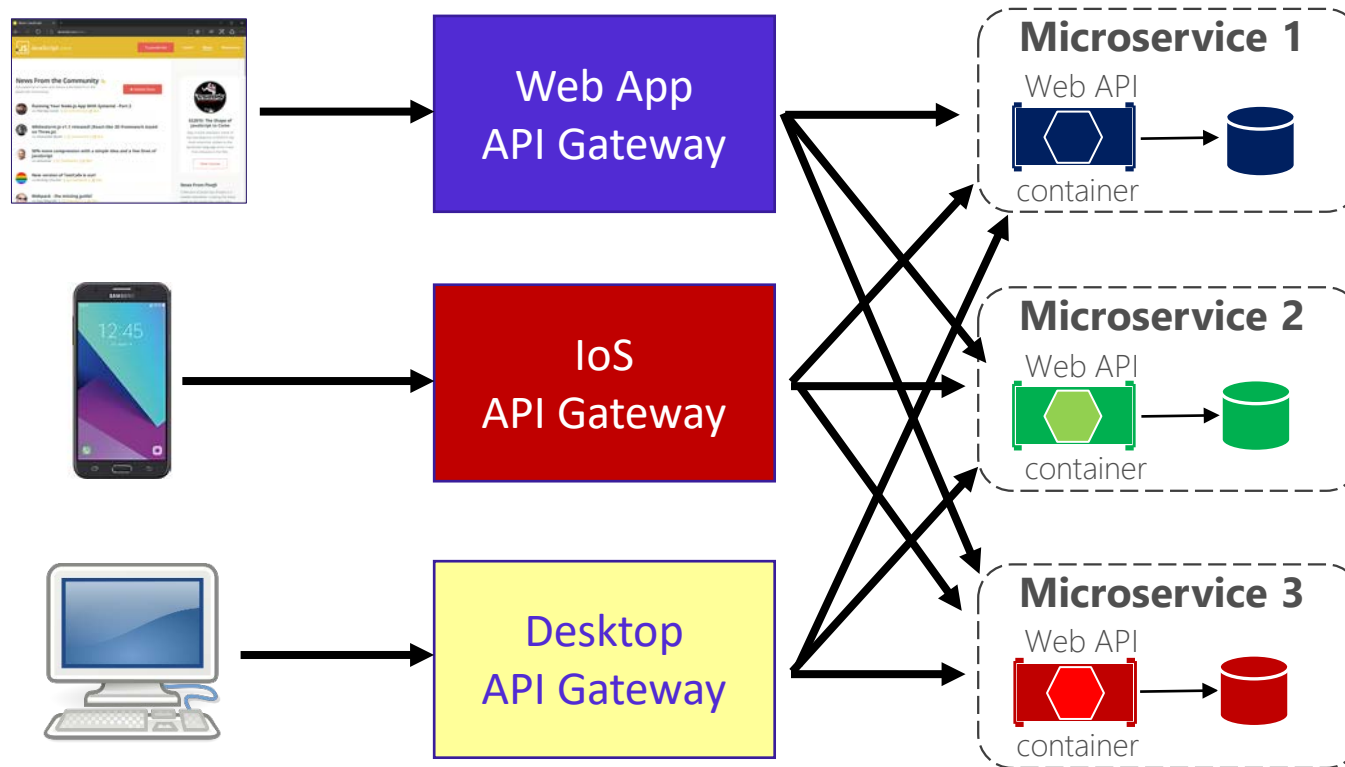
- Front-end service that encapsulates core backend services
  - Exposes single point of entry
  - Routes requests to backend services
  - Insulates client from internal partitioning and refactoring
  - Can fan requests across multiple backend services



- Keep gateway close to backend services to reduce latency
- Ensure availability to prevent single point of failure

# Backends for Frontends Pattern

- Beware of *“overly-ambitious” gateways* – especially in large systems
- Consider multiple gateways
  - Separate by user experience
  - Separate by service category






# API Gateway Offloading Pattern


- Offload some cross-cutting concerns from backend services
  - Centralize in the gateway
  - Simplify backend services

## Centralize Cross-Cutting Concerns...

- Service discovery
- Correlation
- Response caching
- Resiliency logic
- Metering
- Throttling
- SSL termination
- Protocol translation



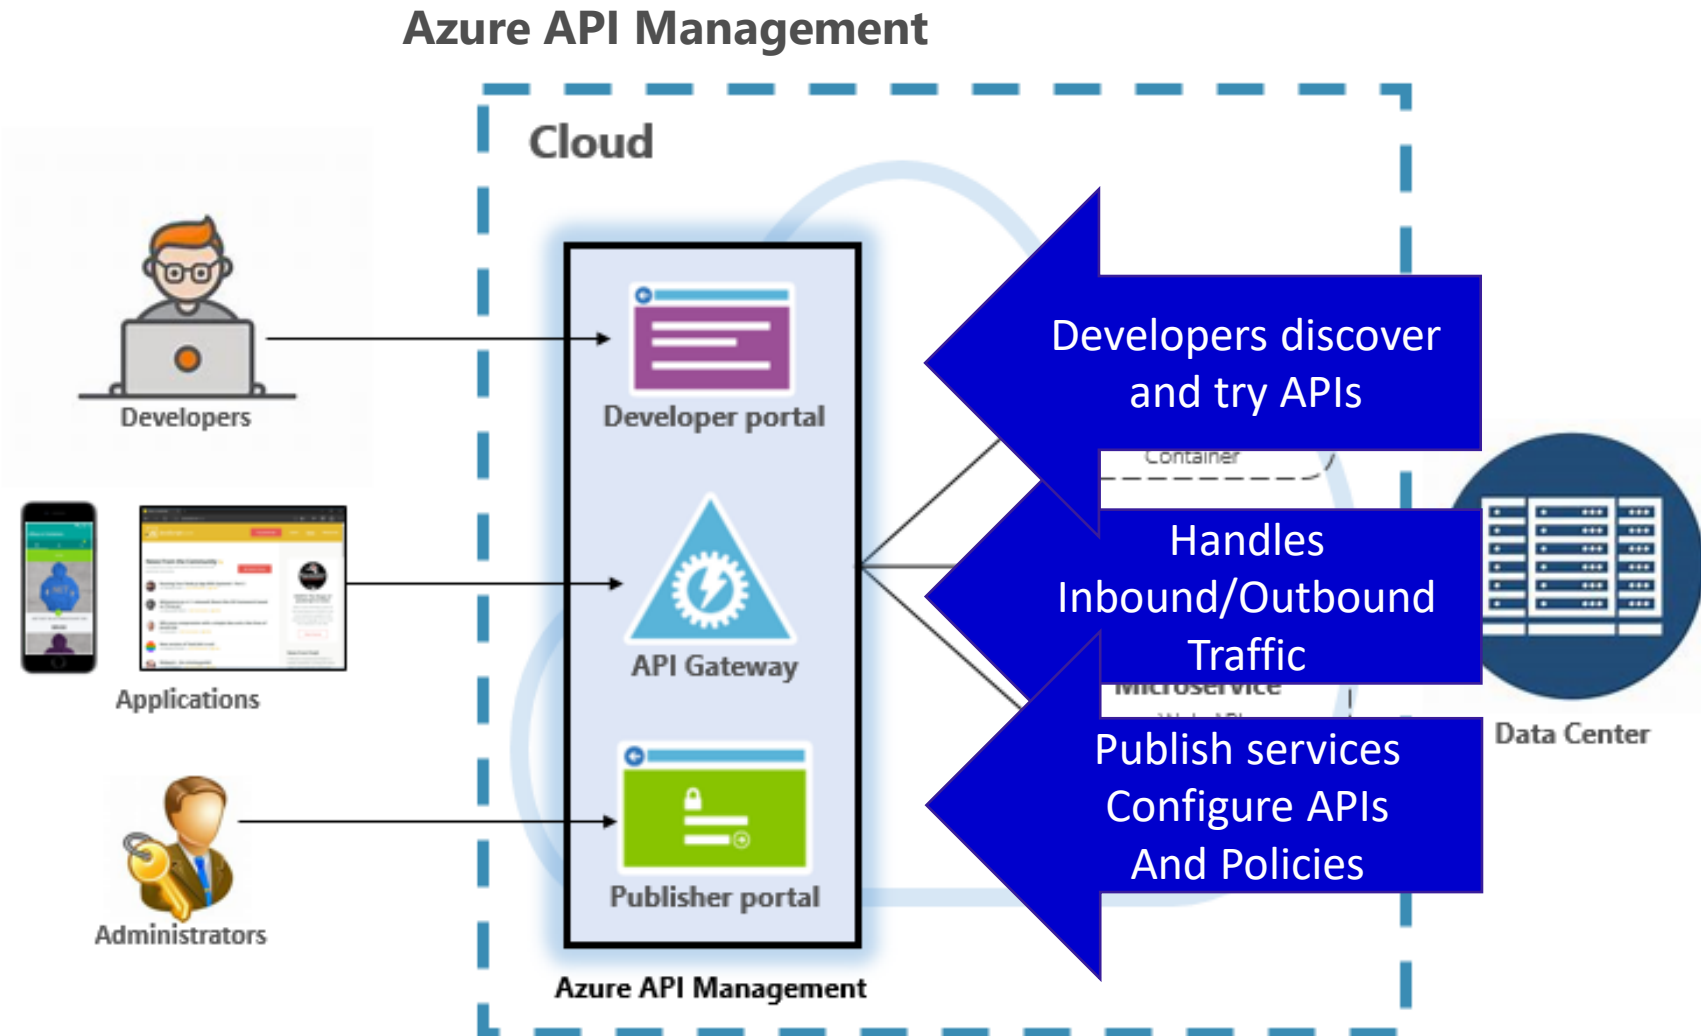
# Demo: API Gateway



# Azure API Management



- Managed service from Azure – Gateway as a Service



# Azure API Management



- Can expose services from Azure, on-prem, and other public clouds
- Developers can apply *policies* to each endpoint to affect behavior...
  - Pre-built functionality can execute for each request
  - Applied to inbound calls, outbound calls, or invoked on error
- Policies include...
  - Restrict access
  - Throttle calls
  - Enable caching
  - Control the flow of a service
  - Transform data formats, such as XML to JSON
  - Custom policies

# Ocelot



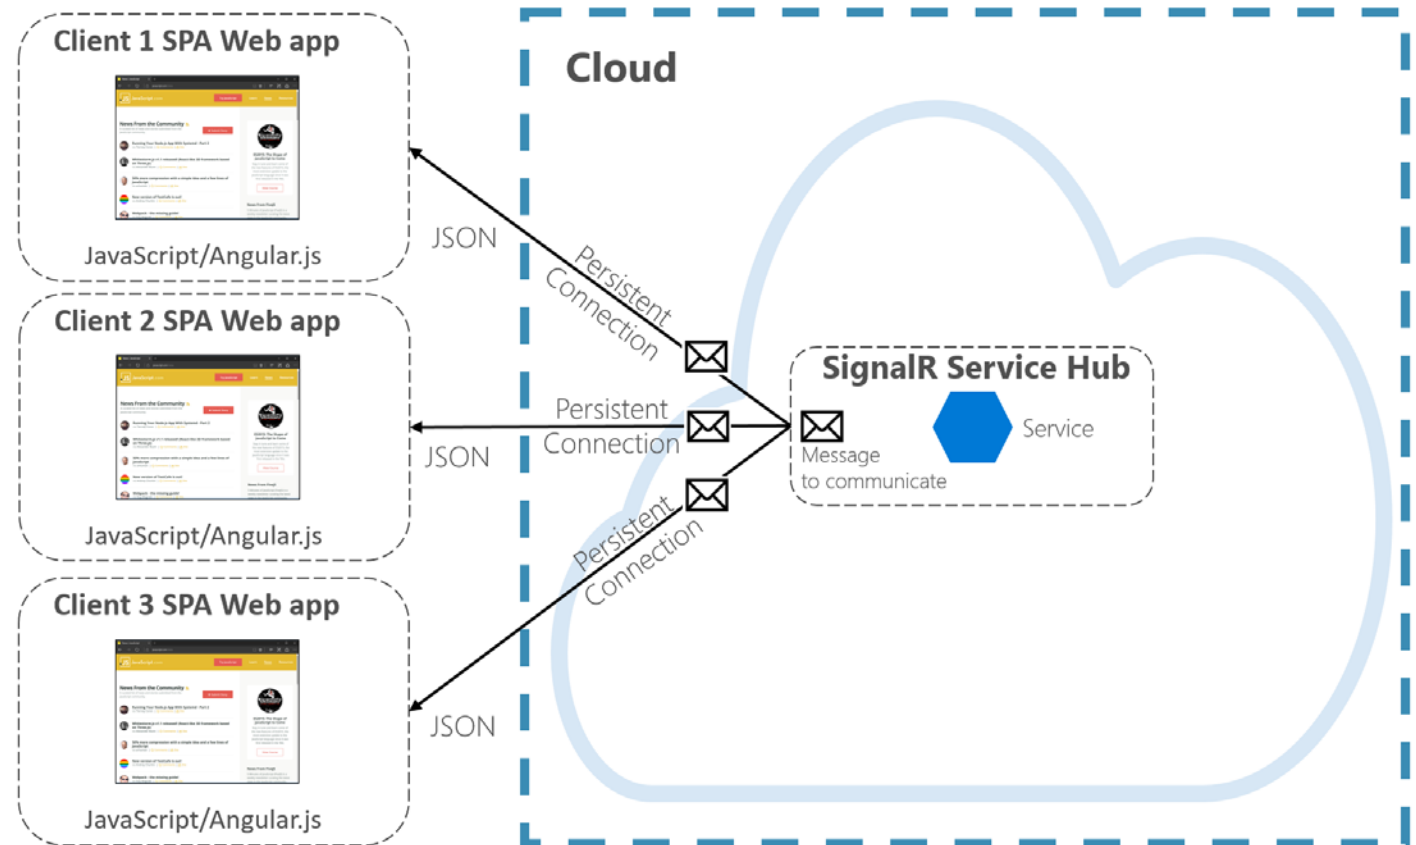
- Open source, .NET Core API Gateway
- Lightweight, extremely simple, scalable - provides numerous features
- Exposes functionality as a set of middleware hosted within an ASP.NET Core
  - Captures incoming Http Request
  - Forwards it through a set of pre-defined middleware, manipulating its state
  - Creates new HttpRequestMessage and forwards to downstream services

Ocelot Features	
Routing	Authentication
Request Aggregation	Authorization
Service Discovery with Consul & Eureka	Throttling
Load Balancing	Logging, Tracing
Correlation Pass-Through	Headers/Query String Transformation
Quality of Service	Custom Middleware


# Real-time communication



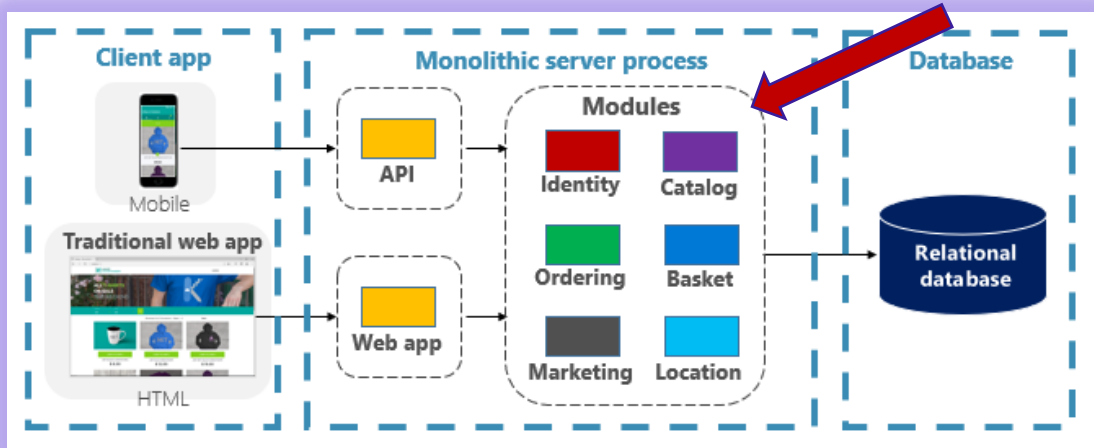
- *Azure SignalR Service* fully managed Azure service that simplifies real-time communication – abstracts the complex plumbing
- Register your client and server can *push data directly* to it
- Clients do not need to poll the server for updates



# Microservice Communication



# Monolithic Communication

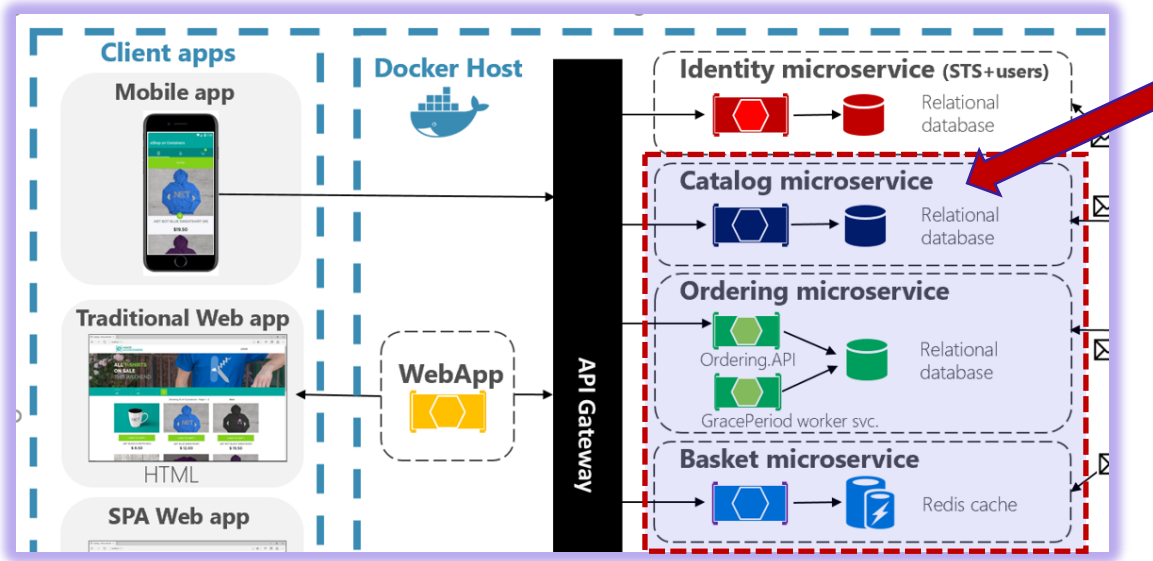


```
// One can directly call another  
var CatalogComponent = new  
CatalogComponent();  
var result =  
CatalogComponent.GetItem(arg1);
```

- For a monolithic app...
  - Communication is straightforward
  - Code modules execute together in server process
  - Can be fast, but results in tightly-coupled code – which can be expensive to maintain, evolve, and scale
- What happens if you *transform* the in-process components to microservices?



# Cloud-Native App



- Life dramatically changes...
  - Calls are now *out-of-process* and communicate across a *network*
    - In-process methods transform to service *endpoints*
    - Each exposes language-agnostic, multi-version *contract*
    - Each must serialize/deserialize arguments/payload (\$\$\$ in memory and CPU)
  - Manage network latency, partitions, transient faults, and unpredictable timing
  - Must authenticate/authorize each service call and encrypt sensitive messages

# How do microservices communicate with each other?

- Two approaches:

Request/Response Model  
(synchronous)  
communication

Publish/Subscribe Model  
(asynchronous)  
communication

- Depending on the *message type...*

# Message Interaction Types

- Message interaction patterns include...

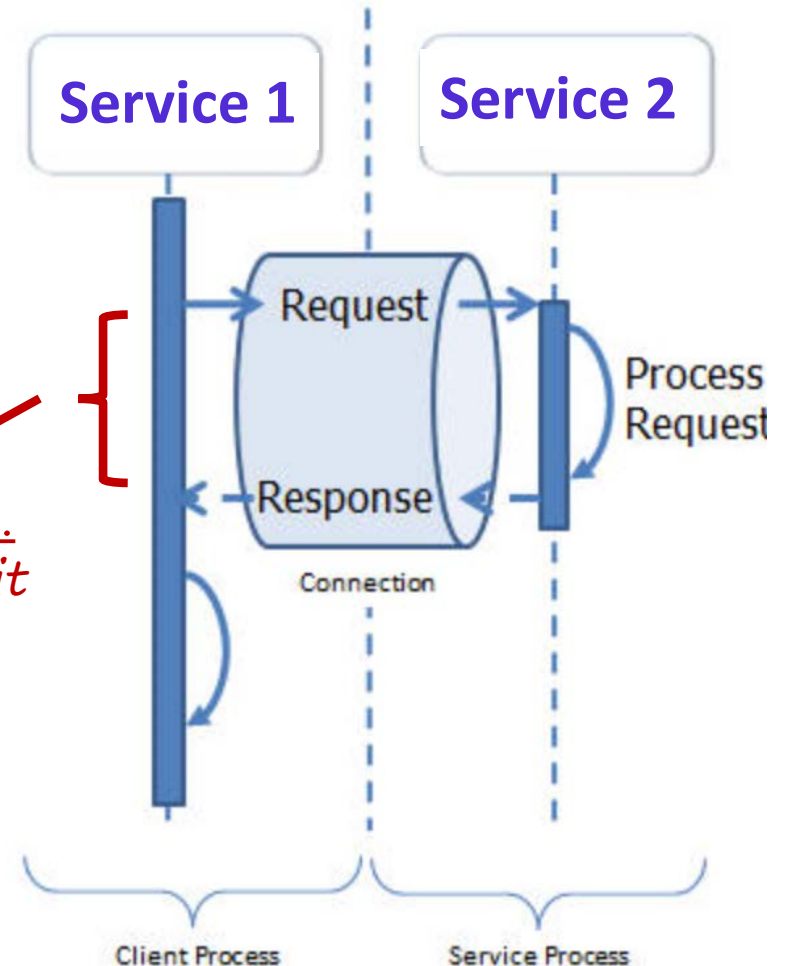
<i><b>Query</b></i>	Client needs response from a service
<i><b>Command</b></i>	Client needs a service to perform an action
<i><b>Event</b></i>	Service reacts to something that's happened in another service

- All three patterns are commonly implemented in cloud applications

# Queries

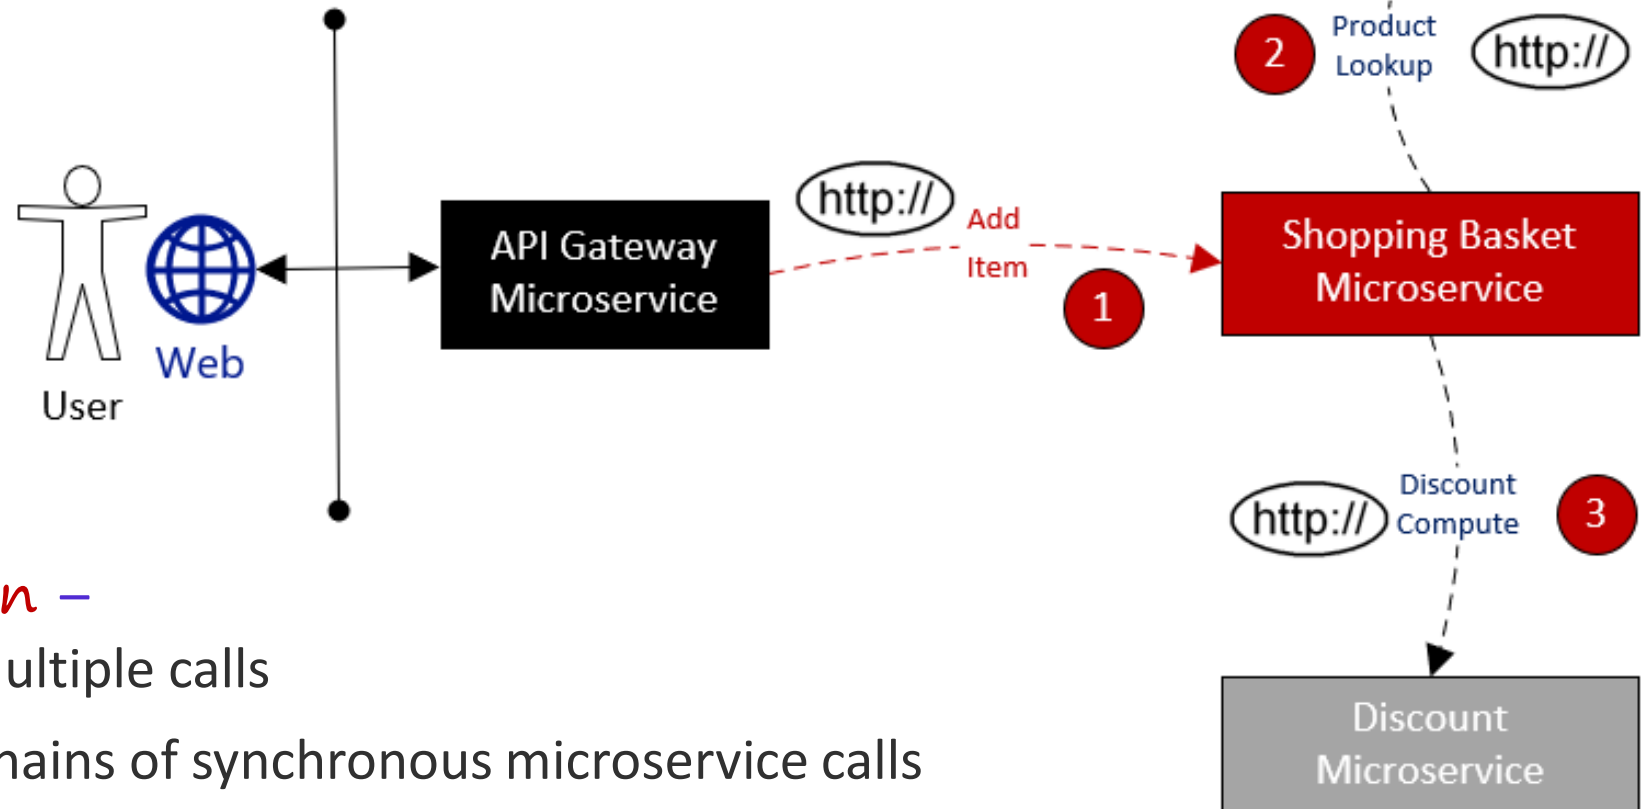
- Calling service has *dependency* upon another service for data or operation
- Needs *immediate response* to complete operation
- Simple to implement
- *Always synchronous*
- Implement async/await pattern to avoid blocking threads

Client process waits...  
Cannot continue until it  
receives a response or  
timeout



# Direct Synchronous Communication

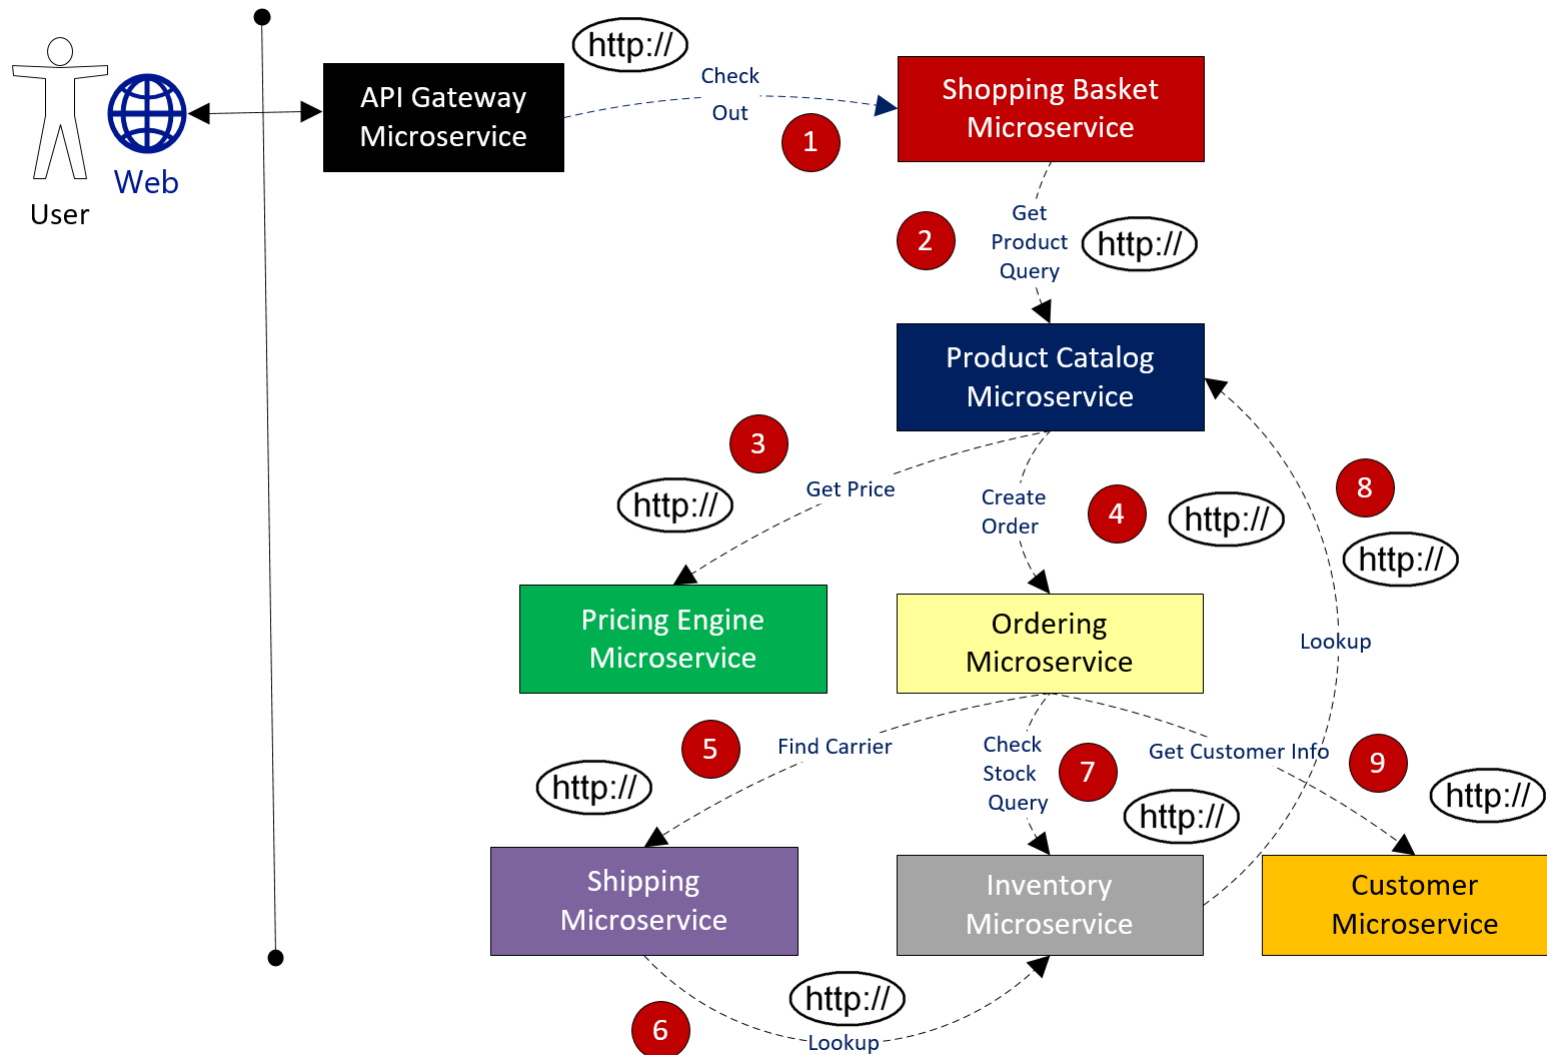
- Couples services, reducing autonomy and architectural benefits
- Impact performance – each call adds latency
- Impact reliability – an unresponsive service can impact entire operation



- *Approach with caution* – especially when invoking multiple calls
- Can lead to long complex chains of synchronous microservice calls

# Chaining HTTP Calls

- Deep chaining or nesting of HTTP calls - simple to implement, but an antipattern

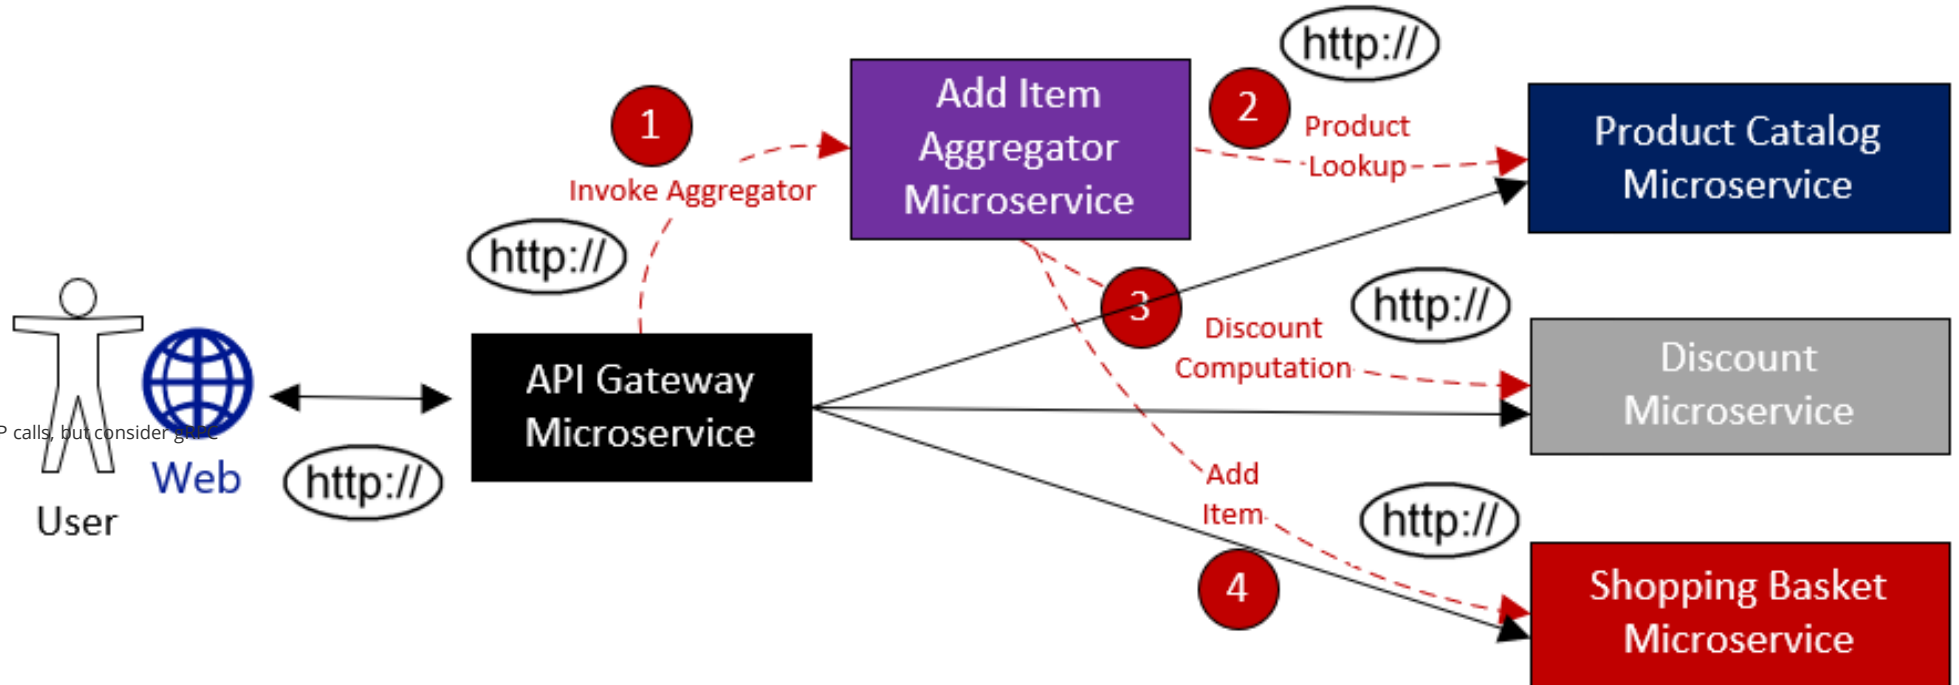


*Not such a good idea!*

# Aggregator Pattern

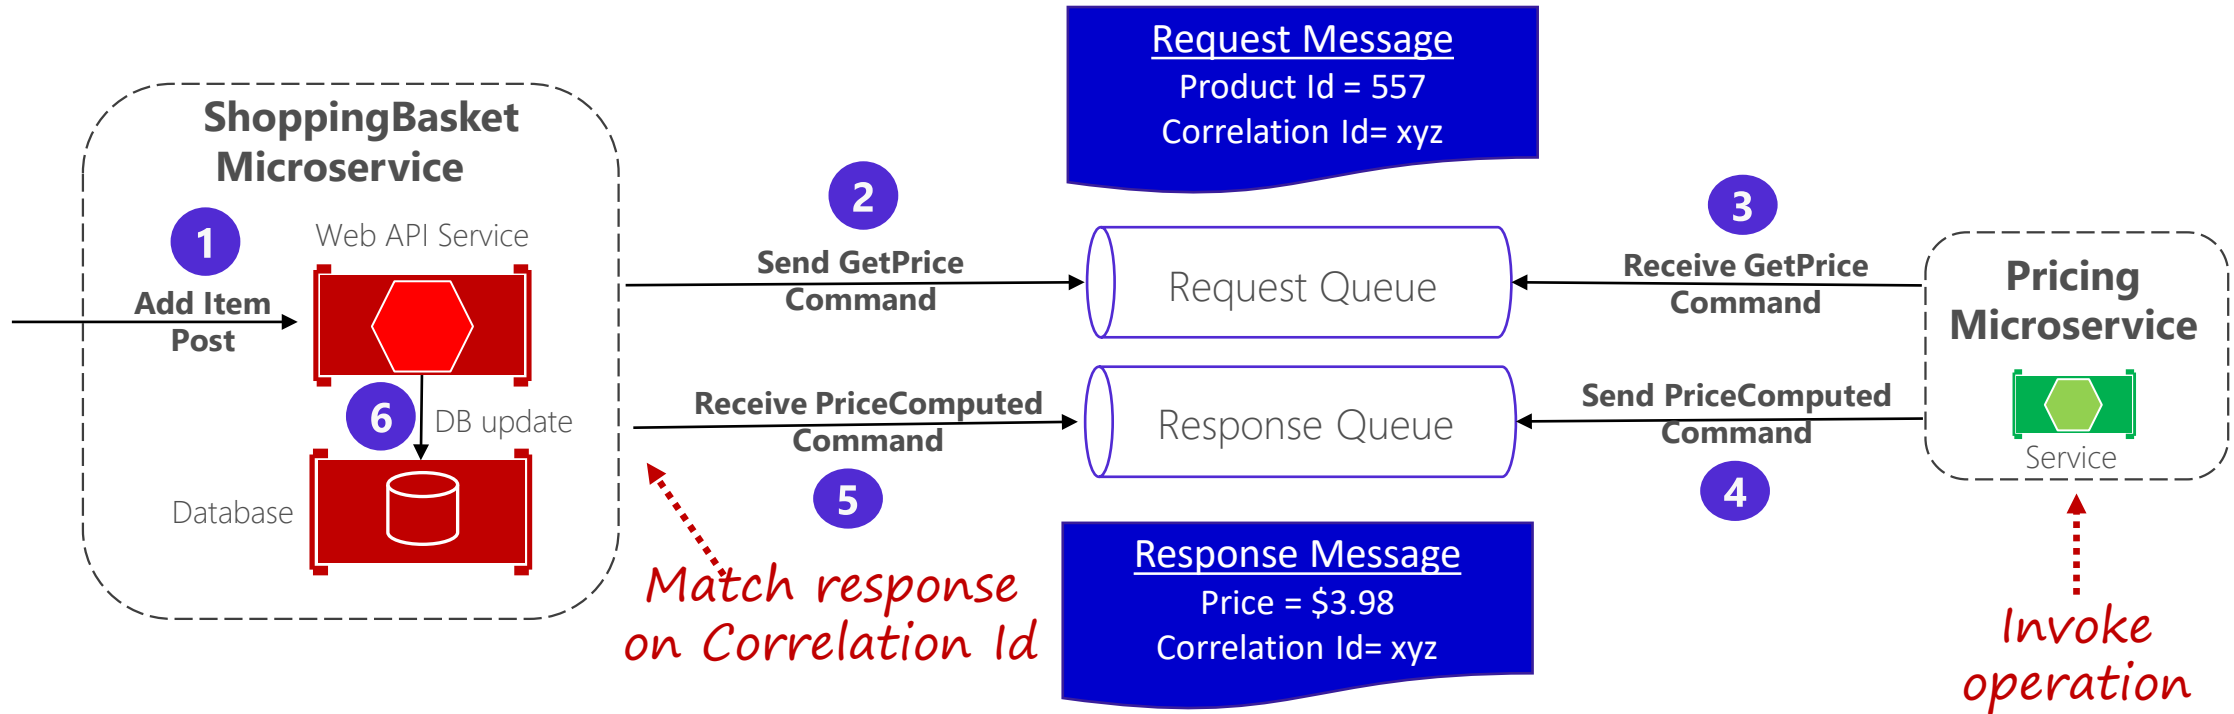
- Aggregator microservice orchestrates a business process
- Orchestrates calls across multiple backend services and aggregates data

- Centralizes an operation
- Can include business logic
- Does not "chain" calls
- Implements synchronous HTTP calls, but consider gRPC



# Request-Reply Pattern

- Referred to as *Sync over Async*
  - Client communication synchronous/backend communication asynchronous

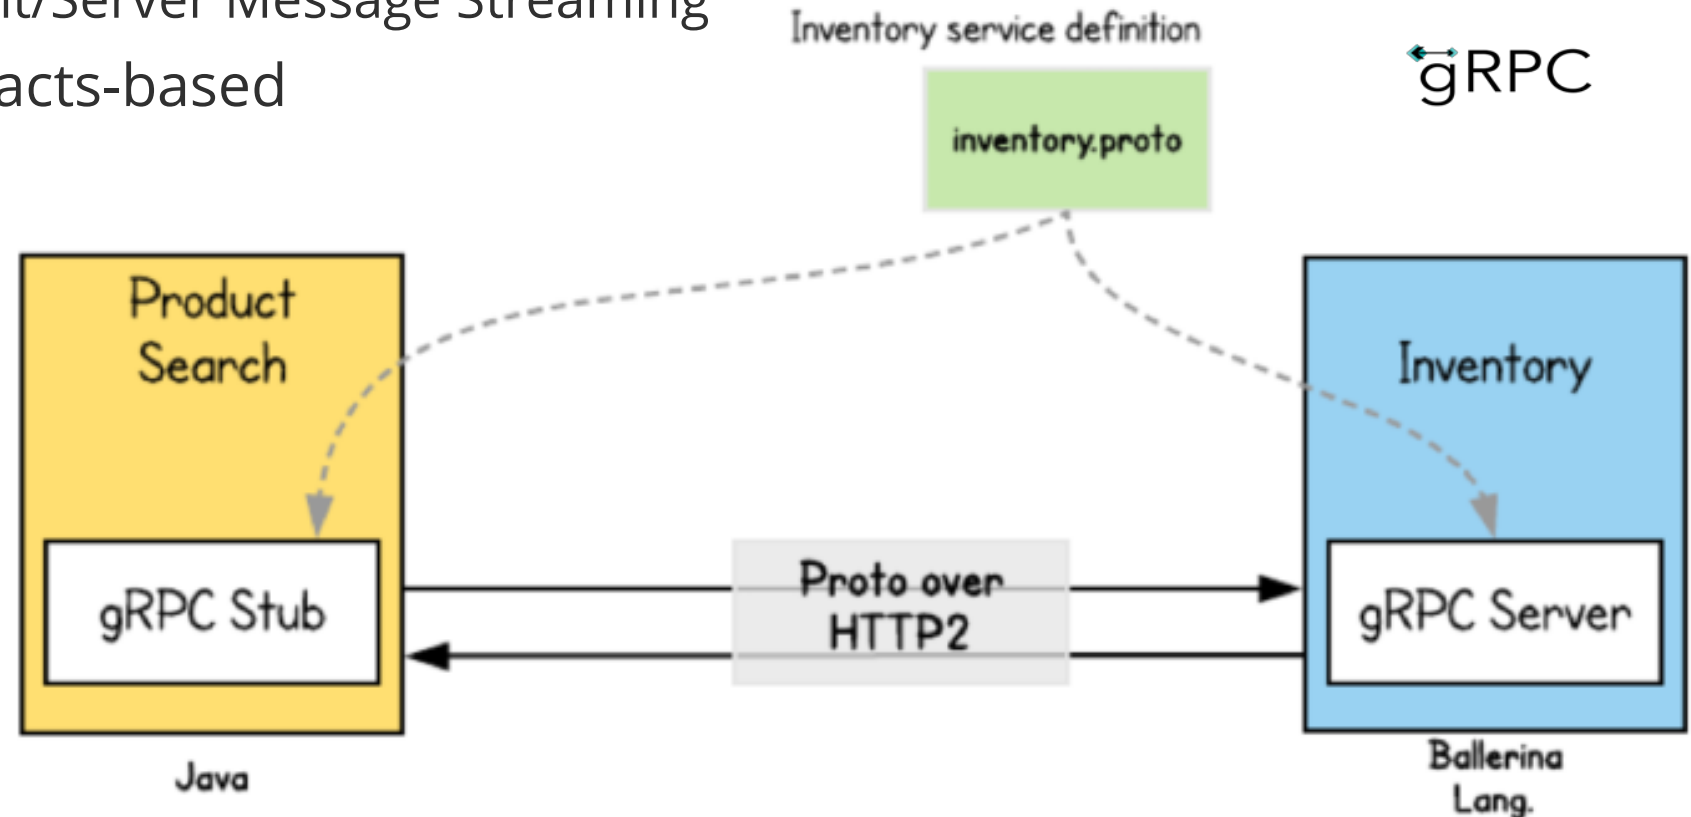


- Decouples calls among backend services
- Adds complexity and not ideal for “awaiting” UI calls
- Consider for systems that implement two-way, real-time communication (SignalR)




# gRPC

- High performance, highly scalable, standards based, open source general purpose RPC Framework
  - gRPC: '**g**' for Google, **RPC** for Remote Procedure Calls
  - Binary Data Representation (compact)
  - Takes advantage of the HTTP/2 feature set
  - Bi-directional Client/Server Message Streaming
- Endpoints are contracts-based

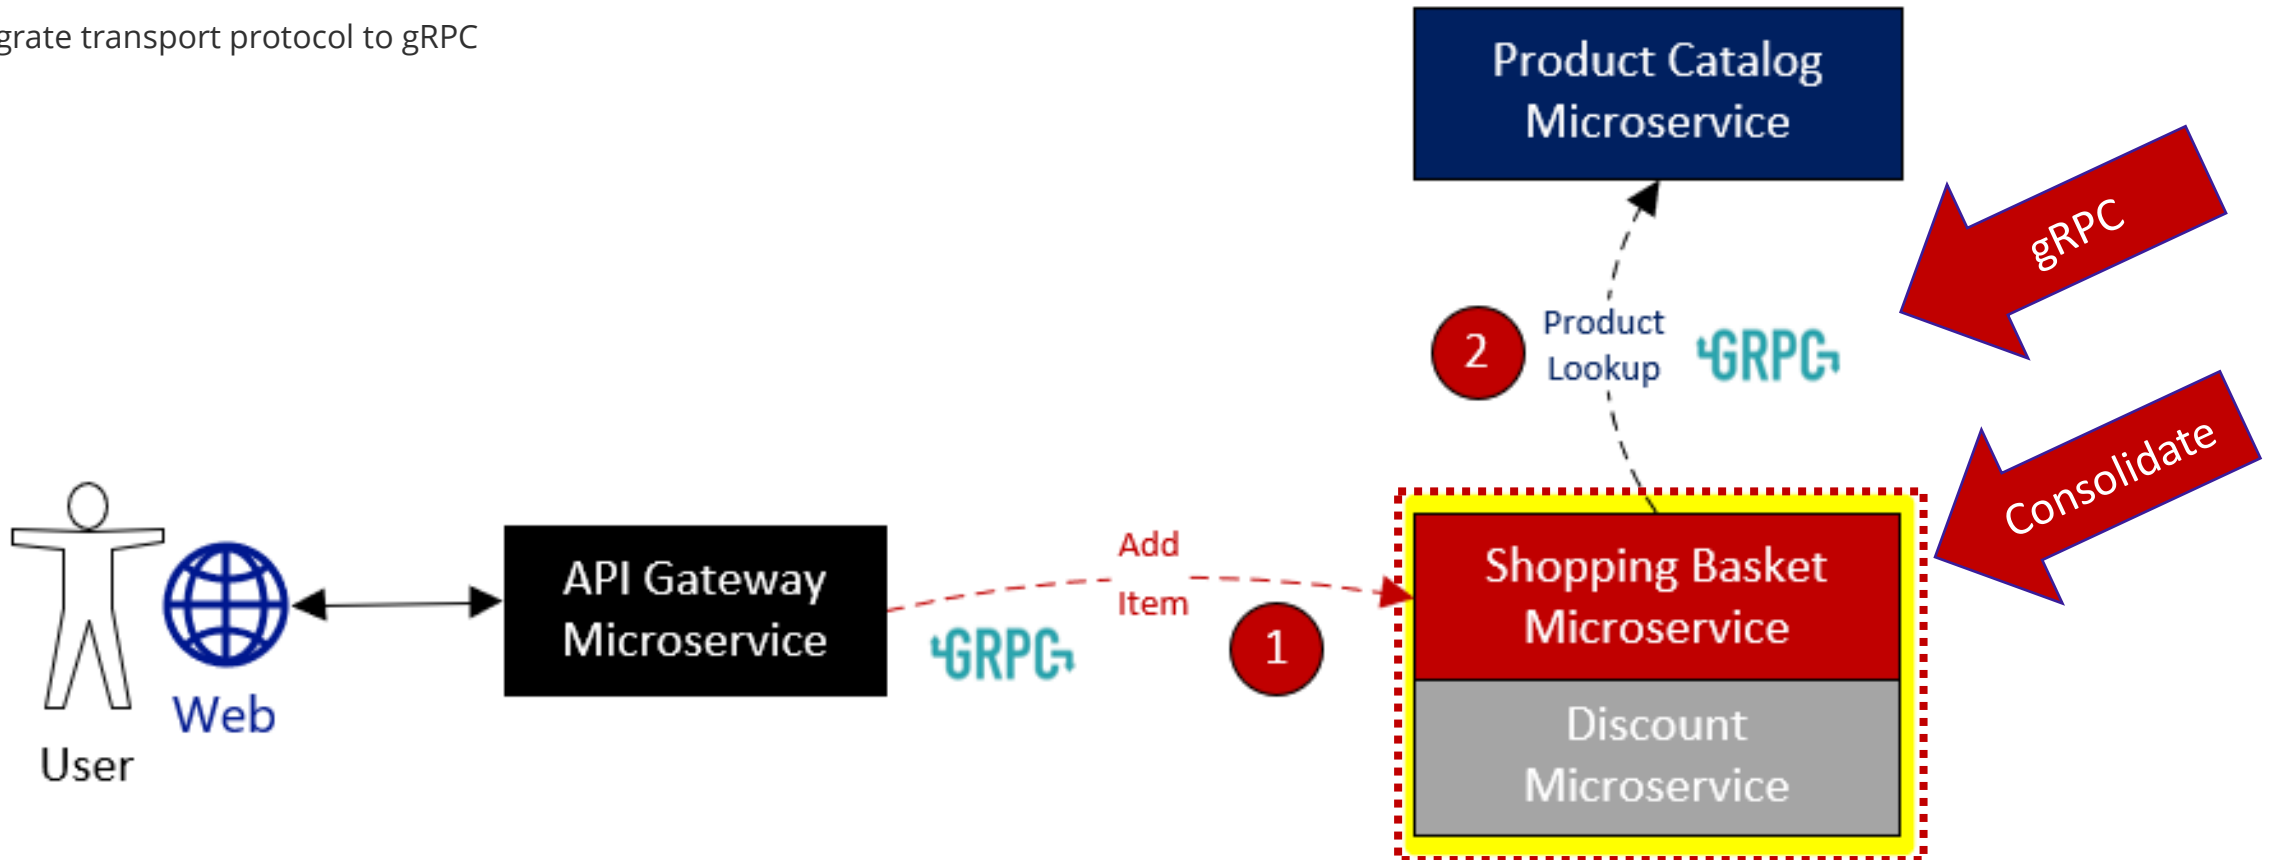


Demo:  
gRPC




# Optimize

- Consolidate Basket and Discount services
  - Tightly integrated
  - Share data
  - High volumes of interservice messages
- Migrate transport protocol to gRPC



Microservice  
Communication:

Commands  
and Events



# Message Interaction Types

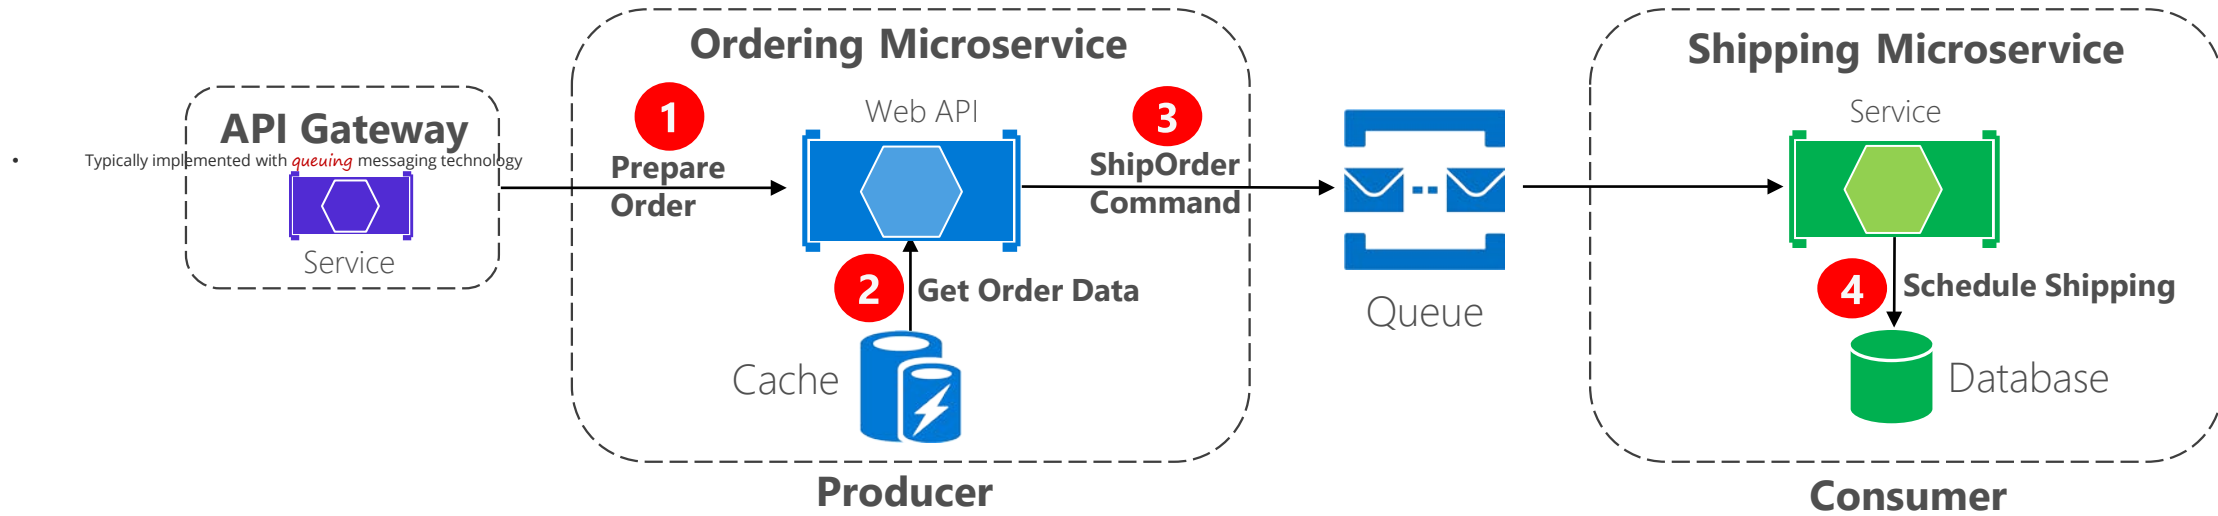
- How do distributed, cloud-based apps/services communicate with each other?
- Message interaction patterns include...

<i>Query</i>	Client needs response from a service
<i>Command</i>	Client needs a service to perform an action
<i>Event</i>	Service reacts to something that's happened in another service

- All three patterns are commonly implemented in cloud applications


# Commands

- Producer/consumer relationship
  - Producer needs consumer to perform an action
  - Producer fires message and forgets
  - Communication moves from synchronous to *asynchronous* exchange
  - Removes cross-service dependency



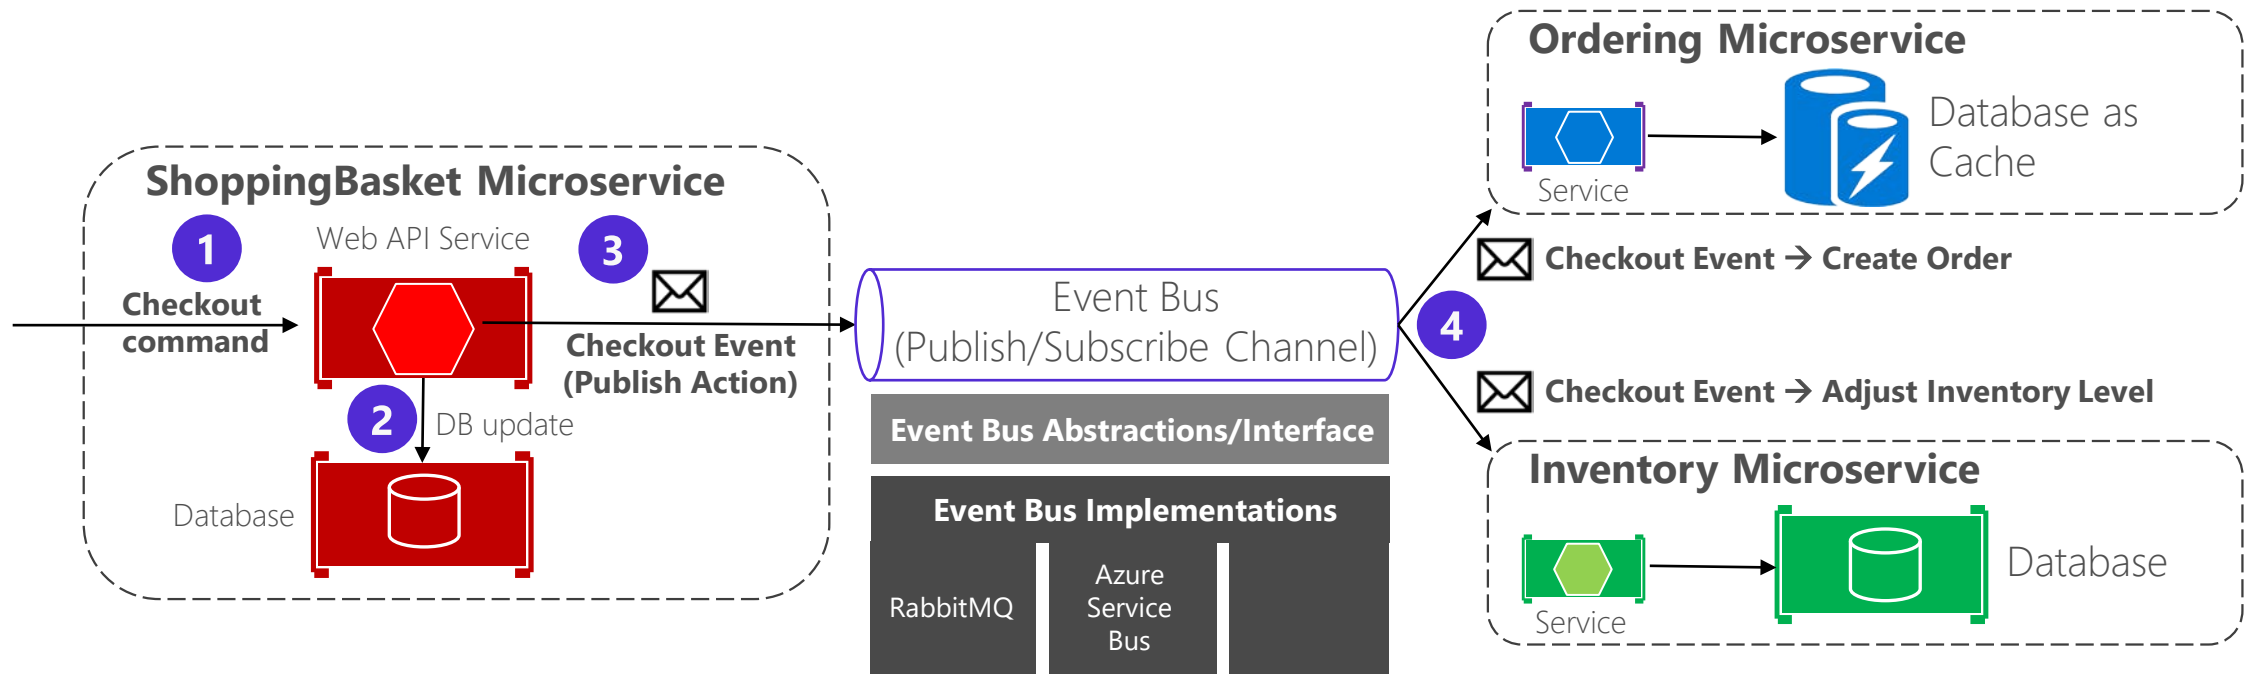
# Queue

- An intermediary construct through which a producer passes a message to a consumer
- Implements a *point-to-point* messaging pattern
- Guarantees each message processed by *at least one consumer*
- Producers and consumers not aware of each other – both have dependency on queue
- Can scale-out one service without affecting the other
- Technologies can be disparate on each side
- Consumer does always not need to run – messages persisted in queue until processed by consumer



# Events

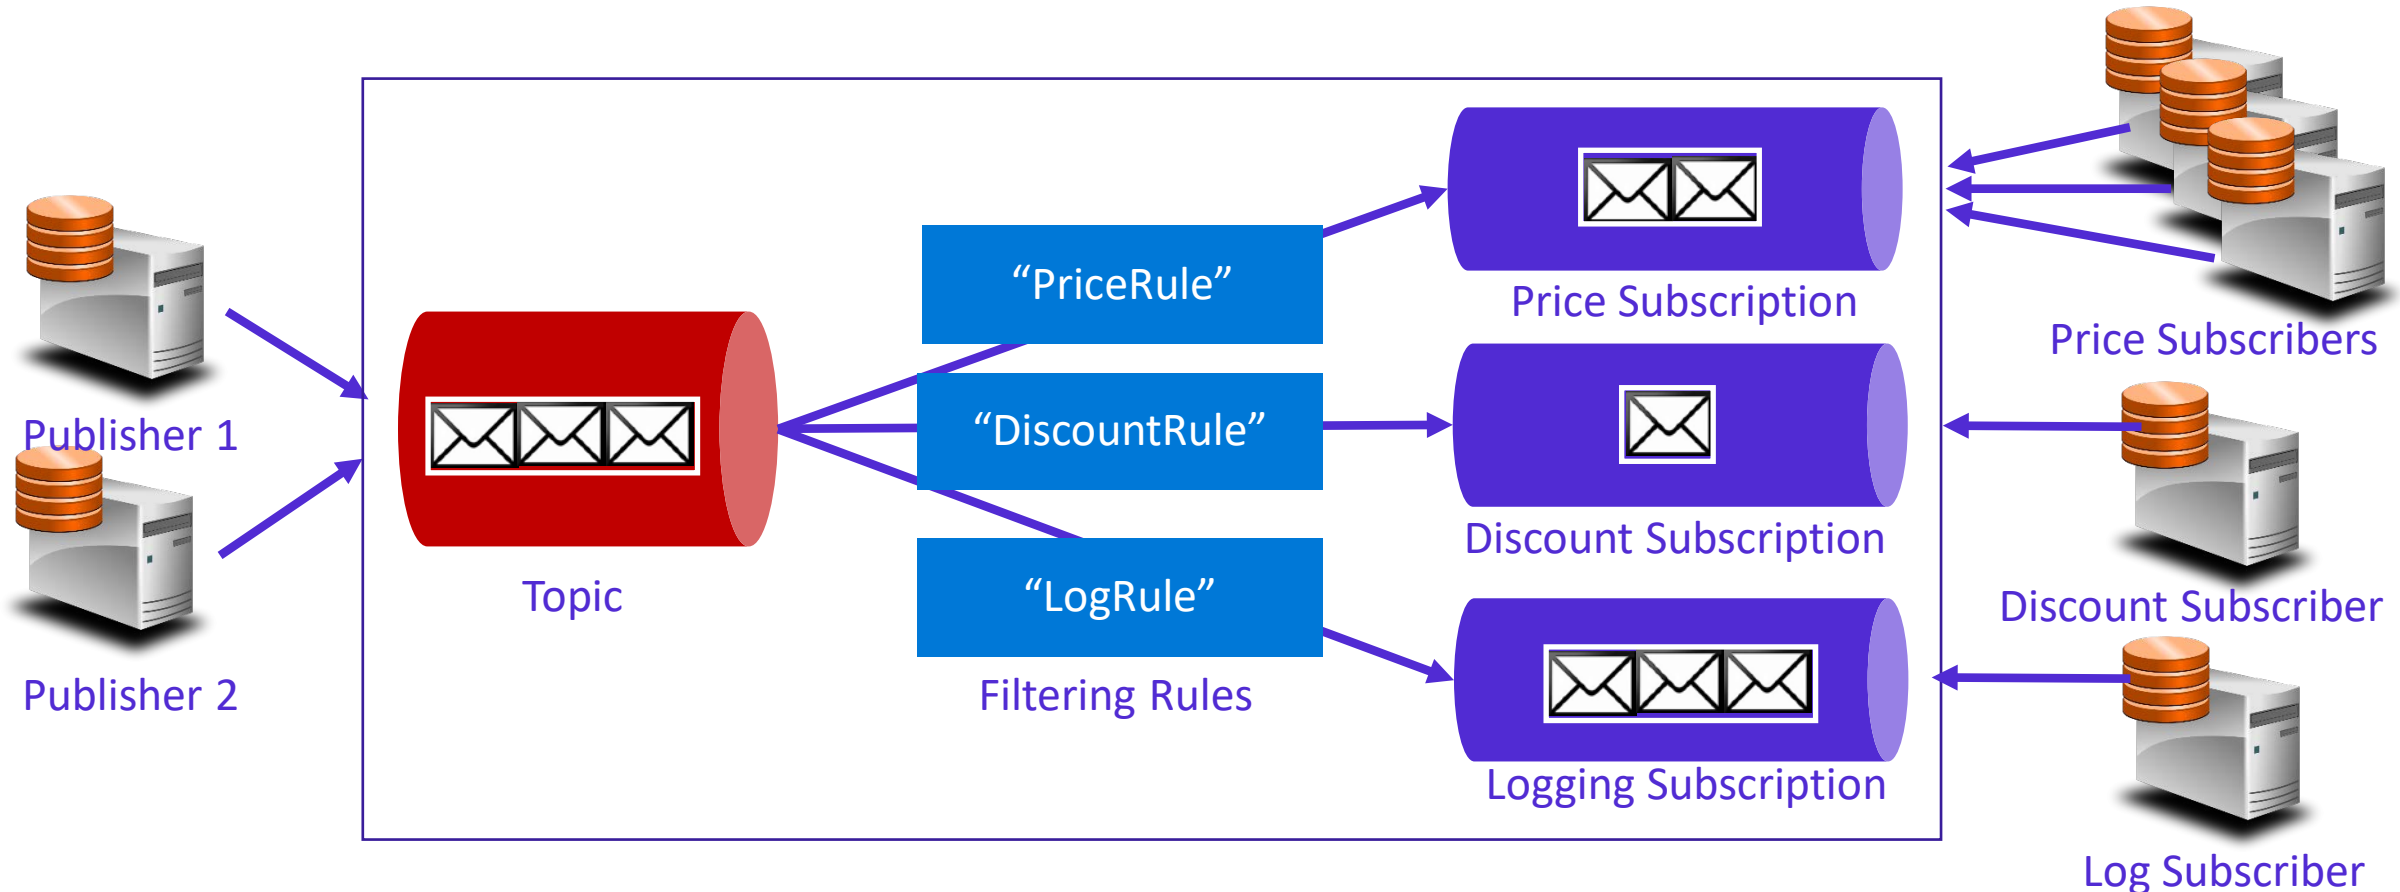
- Publisher/subscriber relationship
  - Publisher raises event upon a state change
  - Subscribers can respond to event - without cross-service dependency
  - Services are unaware of each other
  - Embraces asynchronous communication exchange
- Typically implemented with *topic* messaging technology



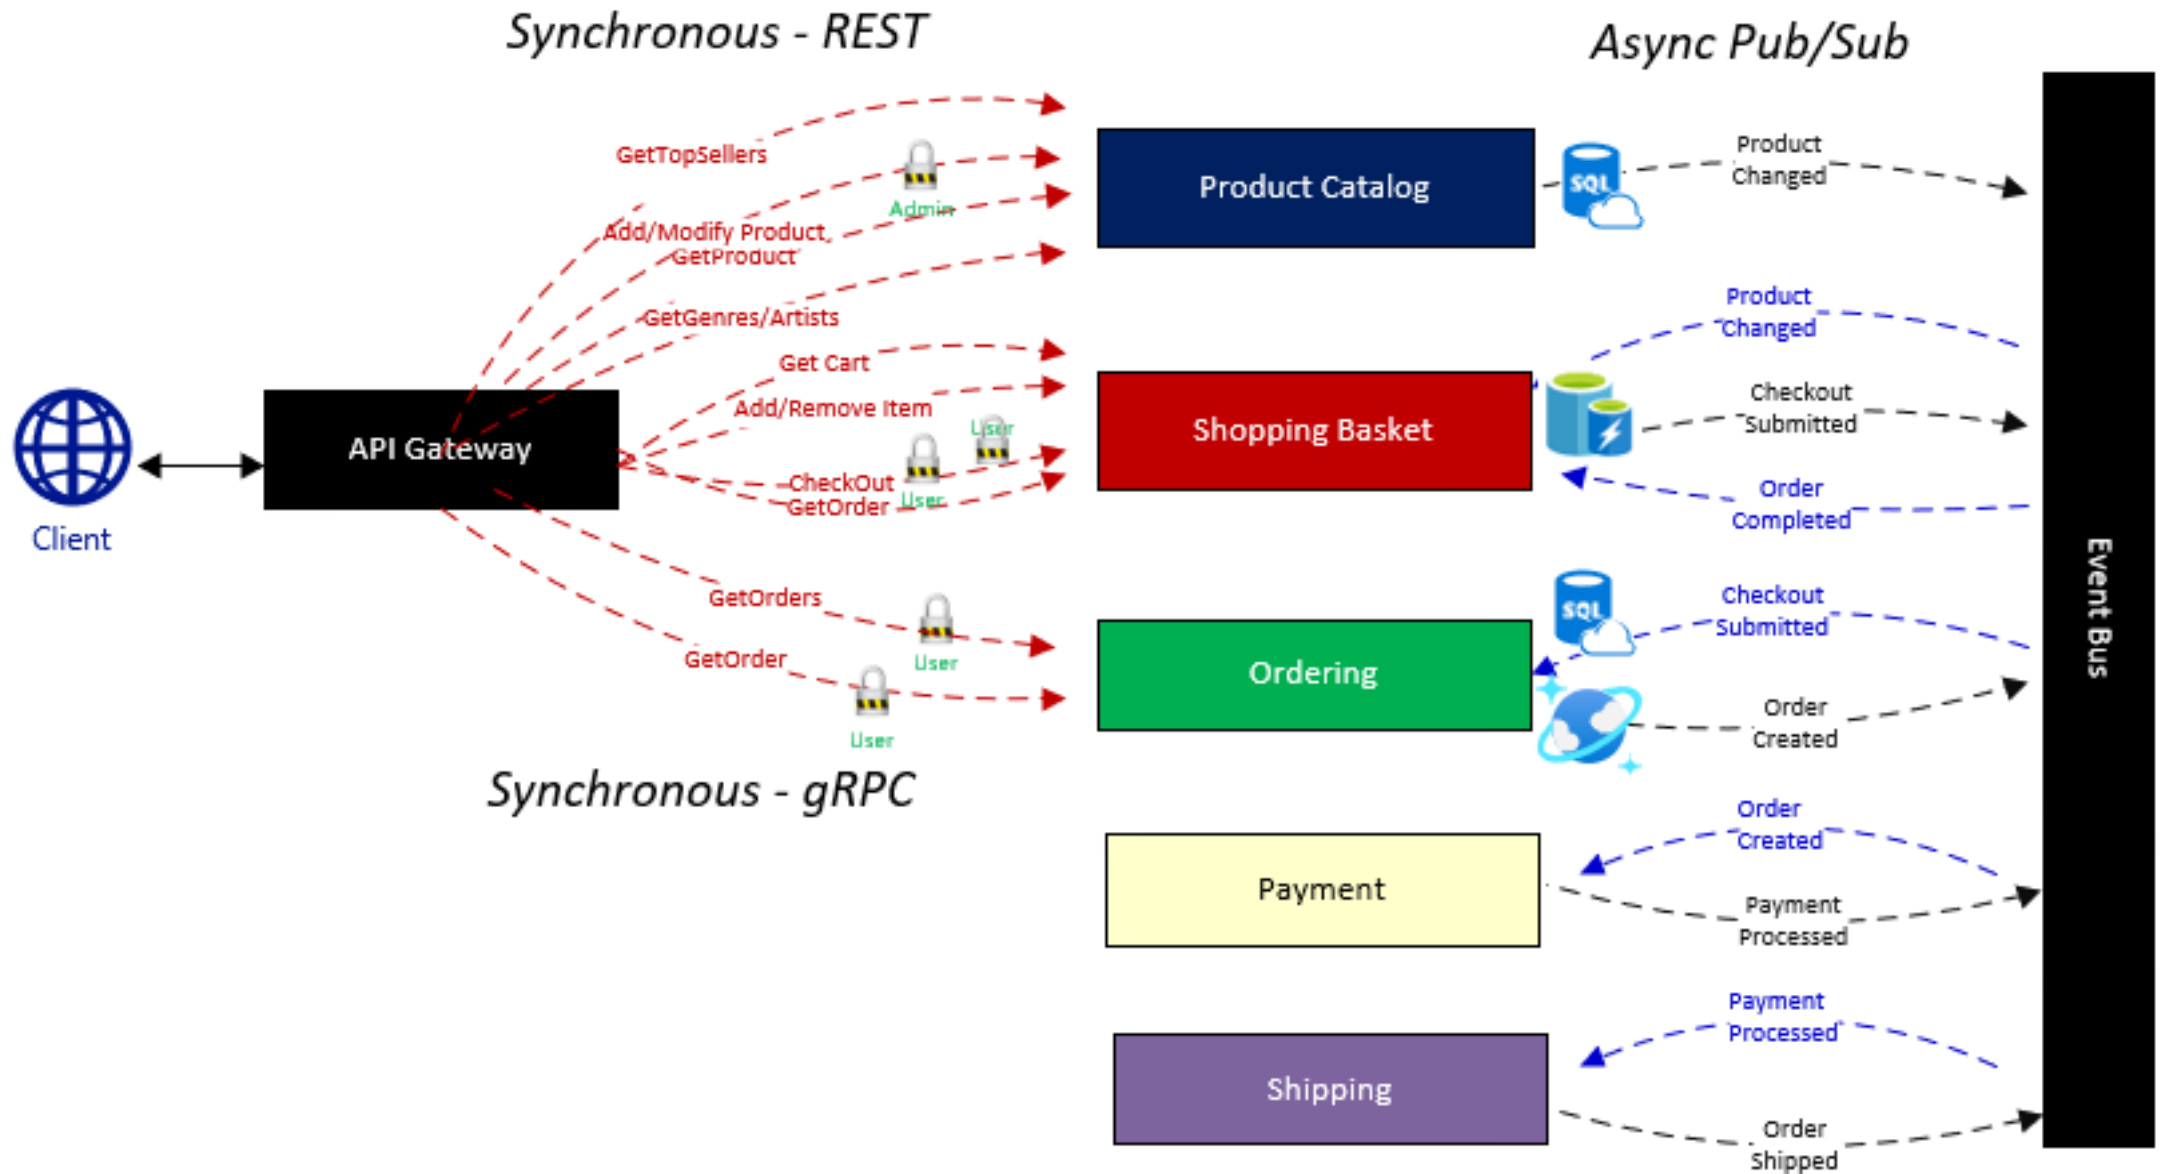


# Topic

- Supports a *one-to-many* messaging pattern to multiple subscribers
  - Publisher service sends message to a topic
  - Subscribing services receive only those message types for which they register
  - Clear separation: Both have dependency on message broker



# Communication in the Reference App



# Demo: EventBus

